


02 • The Inference Stack: Model to GPU to Production

Source: [fanout.sh / inference-eng / in-03](https://fanout.sh/inference-eng/in-03) · ~9 min · Next: *Math for memory, throughput and speedups*


IN SHORT

Between pressing Enter and seeing the first word (~0.5 s), a request crosses about **8 layers** of software and hardware. Each layer exists to answer one question the others can't. The model says **what** to compute; kernels and the GPU decide **how fast**; the engine decides **who runs each step and where their memory goes**; the server and fleet decide **which replica, and how many**. The rest of the course climbs this map.



The stack from user to silicon, grouped by the question each group of layers answers.

1. The model: a file that says what, never how



Llama 3.1 8B: 8B weights in 32 identical layers, plus a recipe and a tokenizer.

- On disk a model is **weights** (a file of numbers) plus a short **recipe** for combining them.

- Llama 3.1 8B: **~8B weights**, **32 identical layers**, 2 bytes each in 16-bit, so **~16 GB**.
- A **tokenizer** sits beside it and converts text to integer IDs and back.

KEY IDEA

The model says nothing about how fast, on which chip, or for how many users at once. Every other layer exists to answer those questions.

2. The GPU: arithmetic waits on memory

- An H100 has two parts that matter: **132 streaming multiprocessors** (thousands of arithmetic units) and **80 GB of HBM** right next to the chip.
- The 16 GB of weights must live in HBM, and **every new token reads every weight once**:

$$t_{\text{step}} = \frac{16 \text{ GB}}{3.35 \text{ TB/s}} \approx 4.8 \text{ ms} \Rightarrow \leq 200 \text{ tok/s per stream}$$


KEY IDEA

This one division, **bytes ÷ bandwidth = time**, is the most useful number in inference. No software above the chip can beat it for a single stream; the arithmetic units mostly wait on memory.

3. Framework, kernels, CUDA: translating math for the chip

Layer	Role	Examples
Framework	Expresses the model as tensor ops (matmul, norm, attention)	PyTorch
Kernels	Small GPU programs, each launched across thousands of threads	cuBLAS (matmul), FlashAttention
CUDA runtime + driver	Hands kernels to the hardware	CUDA

The same math can run at very different speeds depending on the kernel:




Left: fusing three kernels cuts the trips to memory. Right: per-kernel launch costs add up across hundreds of kernels per token.

- **Fusion:** three separate kernels each read from and write back to memory; one fused kernel makes a single round trip.
- **Launch overhead:** each launch costs a few μs of CPU time, repeated across **hundreds of kernels per token**.
- This layer answers: **how fast does this exact math run on this exact chip?**

4. The inference engine: who runs, and where their memory goes

Engines like **vLLM, SGLang and TensorRT-LLM** turn a model into a service. They're built around one trick: if 64 users each need a token, **read the weights once and serve all 64 in the same step**. The step gets only a little slower, but it does $64\times$ the work.



Requests wait in a queue and join the batch at step boundaries; the scheduler re-decides the batch on every step.

- **Job 1, the scheduler:** decides, every single step, which requests join the batch.
- **Job 2, the memory manager:** each request carries a **KV cache** (keys and values of every token it has seen), about **128 KB per token** for Llama 8B. A few thousand tokens $\times$ dozens of requests and the cache competes with the weights for HBM, so the engine hands out and reclaims cache in **small blocks**.


H100 memory: the KV cache fights the weights for the same 80 GB



After 16 GB of weights, the remaining 64 GB of HBM is the KV cache's budget.

5. Server and fleet: which replica, and how many


- **Replica** = one engine on one GPU, fronted by an **API server** that speaks the OpenAI HTTP format, runs the tokenizer and streams tokens back.
- **Router**: a plain load balancer spreads requests evenly. An **inference-aware router** sends a request that shares a long system prompt to the replica that already has it cached, turning hundreds of ms of repeated prefill into almost nothing.



Same 12 requests, two shared system prompts: routing by prefix avoids re-prefilling prompts that are already cached.

- **Autoscaler**: adds or removes replicas. It watches **queue length and KV cache use, not CPU load**, because the CPU is mostly idle. Scaling up is slow: **16 GB of weights must load** before a new replica serves anyone.
- **Metrics and tracing** on every layer tell you which one is hurting.
- The engine optimizes **one replica**; the fleet layer is **everything across replicas**.

6. One request, down and back



Route (~1 ms), tokenize, queue, prefill, then decode at ~5 ms per token, each token streaming back up the stack.

1. The **router** picks a replica in about **1 ms**.
2. The **API server** turns the text into ~**500 token IDs**.
3. The **scheduler** queues the request; at the next step it joins the batch.
4. **Prefill**: all 500 prompt tokens go through the model at once, filling the KV cache; the first token appears (**time to first token**).
5. **Decode**: one token per step, **~5 ms each**, each token climbing back through engine, server and router to the screen.

KEY IDEA

Almost all the time is spent on the GPU, but the layers above decide how well that GPU time is used.

7. Recap

Lesson 02 on one card: four questions, four groups of layers

WHAT

model + PyTorch
weights + recipe: $8B \times 2B = 16GB$
says nothing about speed

HOW FAST

kernels + CUDA + GPU
 $\text{bytes} \div \text{bandwidth} = \text{time}$
fuse kernels, cut launches

WHO & WHERE

inference engine
scheduler picks each step's batch
memory manager hands out KV blocks

WHICH & HOW MANY

API server + router + autoscaler
route by shared prompt
scale on queue and KV use, not CPU

Interview cheat sheet

Revise from these pages alone. ♦ = my addition beyond the video; everything else is from the lesson. Builds on 01's derivations.

Say it in 30 seconds

A request crosses about eight layers. The model defines the math; kernels, CUDA and the GPU set how fast it runs, bounded by bytes ÷ bandwidth. The inference engine batches requests every step and manages KV cache memory. The API server, router and autoscaler pick the replica and the replica count. Nearly all time is GPU time; the upper layers decide how well it's used.

Only numbers worth memorizing

- **H100**: 80 GB HBM, 132 SMs. **Llama 3.1 8B**: 32 layers. Routing ~1 ms; decode ~5 ms/token.

Derive, don't memorize

Every derivation here is the same idea: **time = bytes ÷ bandwidth**, or **capacity = bytes ÷ bytes-per-item**. Per-step time and KV per token: see 01, derivations 1 and 4.

1. KV capacity = memory left after the weights ÷ cache per token ♦

$$N_{\text{tokens}} \approx \frac{\text{HBM} - \text{weights}}{\text{KV/token}} = \frac{80 - 16 \text{ GB}}{128 \text{ KB}} \approx 488\text{k} \approx 119 \text{ requests} \times 4\text{k}$$

INTUITION

That token budget, not compute, caps the batch. It's why engines manage KV in blocks and why quantizing weights or KV frees room for more users.

2. Launch overhead = kernels × launch cost, against the step time ♦

$$\frac{N_{\text{kernels}} \times t_{\text{launch}}}{t_{\text{step}}} = \frac{300 \times 4 \mu\text{s}}{4.8 \text{ ms}} \approx 25\%$$

INTUITION

Illustrative numbers, but the shape is real: at small batch, CPU launch time is a big slice of each step. Fusion (and ♦ CUDA graphs) cut it.

3. Cold start = weight bytes ÷ load bandwidth ♦

$$t_{\text{cold}} \approx \frac{16 \text{ GB}}{2 \text{ GB/s}} = 8 \text{ s}$$

INTUITION

Assumes ~2 GB/s from local disk; from network storage it's often slower. Scaling reacts late, so autoscalers must scale before the queue explodes.

Layer → question → what to watch

Layer	Answers	Tools	Watch / lever
Model + framework	What to compute	Weights, recipe, tokenizer, PyTorch	Model size, precision
Kernels + CUDA + GPU	How fast the math runs	cuBLAS, FlashAttention, CUDA	Bytes moved, fusion, launches
Inference engine	Who runs each step; where memory goes	vLLM, SGLang, TensorRT-LLM	Batch size, KV usage
API server	Speaks HTTP, tokenizes, streams	OpenAI-compatible server	Streaming latency
Router + autoscaler	Which replica; how many	Prefix-aware router	Queue length, cache hits

Symptom → layer to investigate

Symptom	Layer	First fix
GPU idle between kernels, CPU busy	Kernels	Fuse kernels; cut launches
Same math slower than expected	Kernels	Better kernels (cuBLAS, FlashAttention)
Repeated system prompts inflate TTFT	Router	Prefix-aware routing
Queue grows while CPU looks idle	Autoscaler	Scale on queue length and KV use
New replicas arrive too late	Fleet	Derivation 3: scale earlier, load faster

Rapid-fire Q&A

Question	Crisp answer
What does an inference engine add over PyTorch?	A per-step scheduler (batching) and a KV cache memory manager.
Why not autoscale on CPU?	The work is on the GPU; the CPU idles. Watch queue length and KV use.
Why is a round-robin load balancer wasteful?	It scatters shared prompts, so each replica re-prefills them.
What's a replica?	One engine on its GPU(s), behind an API server.
Where does a request's time go?	Mostly GPU (prefill, then ~5 ms per decode step); upper layers decide utilization.

TRAP

- Treating LLM serving like a stateless web app: round-robin routing and CPU-based autoscaling.
- Assuming the model file determines speed: it only defines the math.