

13 • Logits, Sampling & the Next Token


Source: [fanout.sh / inference-eng / 02-04](https://fanout.sh/inference-eng/02-04) • ~7 min

IN SHORT

After 32 layers, the last token's vector becomes text in three steps: **scores** for every vocabulary entry, **probabilities** from the scores, and **one choice**. The **LM head** (after one last RMSNorm) dots the vector with **128,256 rows** to get **logits**. **Softmax** turns them into probabilities, and only their **differences** matter. Then either take the top one (**greedy**) or **sample**, after **temperature** reshapes the distribution and **top-k / top-p / min-p** cut its tail. An engine does this for **every request in the batch, every step**, on the GPU, then checks for a stop, streams, appends and loops.

1. The LM head: one score per token

The vector gets **one last RMSNorm**, then is multiplied by the **language model head**: a matrix with **one row per vocabulary token** (128,256 for Llama 3). Each row is a direction in the same 4,096-d space; the **dot product** with the vector is that token's score, its **logit**. Rows pointing the same way as the vector get high logits.




$128,256 \times 4,096 \approx 525M \text{ weights} \approx 1 \text{ GB in bf16, read on every decode step.}$

INTUITION

Only the **last position's** logits pick the next token, so during prefill engines run the head for **one position**, not the whole prompt (sizes in 10).

2. Softmax: scores to probabilities

Logits are just scores: they can be **negative** and don't add up to anything. Toy vocabulary, prompt "The capital of France is":



Raise e to each logit (all positive, gaps stretched), then divide by the total (200.9) so they sum to 1.


KEY IDEA

Adding the same number to every logit **changes nothing**. Only the **differences** between logits matter.

3. Choosing: greedy or sample

- **Greedy**: always take the most likely token (here, Paris). **Fast and repeatable**, great for facts and code. In open-ended writing it can **lock into loops** ("...and the cat sat on the mat and the cat sat on the mat...").
- **Sample**: draw at random, **weighted by the probabilities**.


Temperature reshapes the probabilities before the draw: divide every logit by T , then softmax again.



$T = 1$ changes nothing; $T = 0.5$ stretches the gaps (Paris 96%); $T = 1.5$ squeezes them (Paris 57%). As $T \rightarrow 0$, sampling becomes greedy, which is why APIs treat temperature 0 as greedy.


4. Cutting the tail

Even at a sensible temperature, the $\sim 128,000$ unlikely tokens **share a little probability**. Draw enough tokens and one of the weird ones eventually gets picked. So we **cut off the tail** and rescale what's left to sum to 1.




Top-k keeps a fixed count; top-p keeps the smallest set whose running sum passes p; min-p keeps any token at least a fraction of the top one's probability.

A **fixed k ignores how confident** the model is. **Top-p adapts**: when the model is confident, the nucleus is tiny; when it's unsure, it grows. **Min-p** is a newer rule with the same spirit.




The usual order:



5. Inside an engine

Serving many requests at once, each step produces **one row of logits per request**, and **every request can bring its own settings** (one greedy, another T 0.8 with top-p 0.95). The engine's **sampler applies each request's settings to its own row, all on the GPU, in one pass**.



Out comes one token ID per request. Stop if it's an end-of-text token, the maximum length, or a stop string. Otherwise turn the ID back into text, stream it, and append it so the next forward pass produces the token after it.

Need the same sampled output twice? Fix the random seed.

6. Recap

Lesson 13 on one card

LM HEAD

final norm, then vector · 128,256 rows
one logit per token (any sign)
525M weights ≈ 1 GB per step

SOFTMAX

e^z , divide by the sum
74 / 12 / 9 / 4 / <1%
only differences matter

CHOOSE

greedy: argmax (facts, code; loops)
sample: z / T , cut the tail
(top-k, top-p, min-p), rescale, draw

ENGINE

per-request settings, one GPU pass
stop check, stream, append, loop
fix the seed to repeat

Interview cheat sheet

Revise from these pages alone. ♦ = my addition beyond the video; everything else is from the lesson. Softmax overflow (subtract the max) is in 05; last-row-only logit sizes in 10; Llama 3.1's default sampling settings in 08.

Say it in 30 seconds

The final-normed vector is dotted with the LM head's 128,256 rows to give logits, one score per token. Softmax exponentiates and normalizes them, and only logit differences matter. Greedy takes the argmax: repeatable, good for facts and code, but it can loop. Sampling divides logits by a temperature, cuts the tail with top-k, top-p or min-p, rescales, and draws; $T \rightarrow 0$ is greedy. An engine applies each request's settings to its own logits row on the GPU in one pass, checks stop conditions, streams, appends, and loops. Fix the seed to reproduce samples.

Only numbers worth memorizing

- Toy logits $5 / 3.2 / 2.9 / 2.1 / 0.5 \rightarrow 74 / 12 / 9 / 4 / <1\%$. Paris at $T = 0.5$: **96%**; at $T = 1.5$: **57%**.

Derive, don't memorize

1. Only gaps over T matter

Divide any two probabilities: the softmax denominator cancels, leaving the logit gap scaled by T.

$$\frac{p_i}{p_j} = e^{(z_i - z_j)/T} \quad \frac{p_{\text{Paris}}}{p_{\text{the}}} = e^{1.8} \approx 6 \quad (T=1), \quad e^{3.6} \approx 37 \quad (T=0.5)$$

INTUITION

Add c to every logit and the gap is unchanged, so nothing moves. As $T \rightarrow 0$ every ratio blows up and the top token takes everything: greedy. Temperature never reorders tokens.

2. Cut, then rescale

Keep a set S of tokens, then divide each survivor by the kept mass.

$$p'_i = \frac{p_i}{\sum_{j \in S} p_j} \quad \text{top } p = 0.9 : 74, 86, 95 \Rightarrow 3 \text{ kept}, \quad p'_{\text{Paris}} = \frac{74}{95} \approx 78\%$$

INTUITION

Min-p 0.1 keeps $p \geq 0.1 \times 74\% = 7.4\%$: the same three tokens here. Top-k = 3 keeps them by count.

3. The head is a fixed cost per step ♦

$$V \cdot h \cdot 2 \text{ B} = 128256 \times 4096 \times 2 \approx 1.05 \text{ GB} \quad \frac{1.05 \text{ GB}}{3.35 \text{ TB/s}} \approx 0.31 \text{ ms per decode step}$$

Decoding rules compared

Rule	What it does	Adapts to confidence?	Toy result
Greedy ($T = 0$)	argmax	n/a	Paris, always
Temperature	z / T before softmax	Reshapes, keeps order	Paris 96% (0.5), 57% (1.5)
Top-k	Keep the k most likely	No: fixed count	$k = 3 \rightarrow 78 / 13 / 10$
Top-p	Smallest set with $\Sigma \geq p$	Yes: nucleus grows when unsure	$p = 0.9 \rightarrow 3$ tokens
Min-p	Keep $p \geq m \times p_{\text{max}}$	Yes: relative to the top	$m = 0.1 \rightarrow$ cutoff 7.4%

Symptom $\rightarrow$ diagnosis

Symptom	Diagnosis
Output repeats the same phrase	Greedy on open-ended text: sample, or add penalties
Rare bizarre token mid-answer	No tail cut: add top-p or min-p
Can't reproduce a sampled answer	Seed not fixed

Rapid-fire Q&A

Question	Crisp answer
What is a logit?	The dot product of the final vector with one LM-head row: an unnormalized score.
Why compute logits only at the last position?	Only it picks the next token; prefill runs the head for one position.
Why is temperature 0 allowed?	APIs treat it as greedy, the $T \rightarrow 0$ limit.
Why top-p over top-k?	A fixed k ignores confidence; top-p's nucleus shrinks or grows.
What order does a sampler apply?	Penalties, temperature, cut the tail, rescale, draw.
How do mixed settings share a batch?	Each request's settings apply to its own logits row, in one GPU pass.
When does a request stop?	End-of-text token, maximum length, or a stop string.

TRAP

- Thinking temperature reorders tokens: it only sharpens or flattens.
- Forgetting to rescale after cutting the tail: survivors must sum to 1.
- Treating logits as probabilities: they can be negative and don't sum to 1.