

17 • Build a KV Cache


Source: [fanout.sh / inference-eng / 03-04](https://fanout.sh/inference-eng/03-04) • ~7 min

IN SHORT

A KV cache is **one key tensor and one value tensor per layer**, shaped **[batch, 8 KV heads, max length, 128]** for Llama 3 8B, plus a **pointer** saying how many slots are filled: **128 KiB per token**, so 8,192 reserved tokens is exactly **1 GiB per request**. **Preallocate** the buffer and write each new K/V **in place** (concatenating copies everything every step: 33.6M token copies over 8k tokens). Prefill writes slots 0 ... P-1 in one shot; each decode step writes at the pointer, reads only slots 0 ... pointer, and advances it. Four classic bugs: **wrong RoPE position, reading empty slots, not resetting, copying KV heads**. Prove it with identical greedy tokens, matching byte counts and flat step time. The hidden cost: reserved but **unused memory**.

1. What we store

In every layer, attention produces a key and a value **per token, per KV head**. So each layer gets two tensors:




32 layers × (K, V), each [B, 8, L_max, 128], plus one small but essential piece: the position pointer.

The per-token size is lesson 05's 128 KiB ($2 \times 32 \times 8 \times 128 \times 2$ bytes). Reserve 8,192 tokens and a single request's cache is **exactly 1 GiB**.

2. How the cache grows

```
# naive: a brand-new tensor every step, copying everything
k = torch.cat([k, k_new], dim=2)

# better: allocate once, write into the next empty slot
k_cache = torch.empty(B, 8, L_max, 128)
k_cache[:, :, pos] = k_new
pos += 1
```




Concatenation: more than 33 million token copies just to add 8,192 tokens. Preallocation: one small write per step, whatever the length.

INTUITION


With preallocation the **tensor shapes never change**, exactly what **CUDA graphs and compilers** need. Production code (e.g. torchtune) uses a fixed buffer plus a position counter.

3. Prefill writes, decode appends




- **Prefill:** the whole prompt arrives at once; its K/V go into slots **0 ... P-1 in one shot**, in every layer. Inside the prompt, attention still uses the **causal mask**. The pointer jumps to the end of the prompt.
- **Decode:** each step writes the new token's K/V **at the pointer**, its query reads slots **0 ... pointer and nothing beyond**, then the pointer moves forward by one.

4. Four bugs that catch almost everyone




- Position encoding.** Llama rotates each query and key by an angle set by the token's **position** (RoPE). In decode you feed one token, so it's tempting to call it position 0, but it's really **position 512** (or wherever the pointer is). Get it wrong and the first token is fine, then **gibberish**.
- Reading empty slots.** A preallocated buffer is full of zeros. A zero key gives score **0**, and $e^0 = 1$, **not 0**, so unfilled slots quietly **steal attention weight**. Always **slice or mask** to the filled length.
- Forgetting to reset.** If the pointer isn't cleared between requests, the next user's prompt **attends to the previous user's text**.
- Heads.** Llama stores only **8 KV heads**, each shared by 4 query heads. **Broadcast** them at attention time; don't copy them into the cache.

5. Prove it works



- **Correctness:** generate 100 tokens greedily with and without the cache; the token sequences must **match exactly**. Caching changes **how** we compute, not **what**. Logits won't be bit-identical in 16-bit, but should agree **within a small tolerance**.
- **Memory:** allocated bytes must equal $2 \times \text{layers} \times \text{KV heads} \times \text{head dim} \times \text{length} \times \text{bytes per number}$.
- **Speed:** time each decode step: **flat with the cache**, climbing without.

6. The hidden cost of preallocating



Eight requests reserved for 8,192 tokens each set aside 8 GiB. If they average 1,000 tokens, about 7 GiB sits empty, which limits how many requests fit on the GPU. Paged memory fixes this.

7. Recap

Lesson 17 on one card

SHAPE

per layer: K and V [B, 8, L_max, 128]
× 32 layers + a position pointer
128 KiB/token, 1 GiB per 8k request

GROW

preallocate once, write at pos
not torch.cat (33.6M copies)
fixed shapes: CUDA graphs, compile

FOUR BUGS

RoPE at the true position
mask or slice empty slots ($e^0 = 1$)
reset per request; store 8 KV heads

PROVE IT

greedy tokens identical with/without
logits within tolerance; bytes match
step time flat. Cost: reserved waste

Interview cheat sheet

Revise from these pages alone. ♦ = my addition beyond the video; everything else is from the lesson. Why the cache exists is in 14; KV bytes per token in 05, derivation 2; GQA in 07.

Say it in 30 seconds

Per layer I keep a key and a value tensor shaped $[batch, KV\ heads, max\ length, head\ dim]$, plus a position pointer. That's 128 KiB per token for Llama 3 8B, 1 GiB for an 8k reservation. I preallocate and write in place instead of concatenating, which avoids quadratic copying and keeps shapes static for CUDA graphs. Prefill fills slots 0 to $P-1$ at once under a causal mask; each decode step writes at the pointer; reads only up to it, and advances. The classic bugs are RoPE at the wrong position, attending to zero-filled slots ($e^0 = 1$), not resetting between requests, and copying KV heads instead of broadcasting. I test that greedy outputs match the uncached loop, bytes match the formula, and step time stays flat.

Only numbers worth memorizing

- **Llama 3 8B cache:** per layer $K, V = [B, 8, L_{max}, 128]$; **8,192 tokens = 1 GiB** per request.

Derive, don't memorize

1. Bytes the cache should allocate (test 2)

$$\text{bytes} = 2 \cdot L \cdot H_{kv} \cdot d \cdot L_{max} \cdot b \cdot B = 128\text{ KiB} \times 8192 = 1\text{ GiB per request}$$

2. Why concatenation is quadratic

Each step copies everything so far, then adds one.

$$\sum_{t=1}^L t = \frac{L(L+1)}{2} = \frac{8192 \times 8193}{2} \approx 33.6\text{M copies} \quad \text{vs} \quad 8192\text{ writes}$$

3. Reserved but empty

$$\text{waste} = R \cdot (L_{max} - \bar{L}) \cdot 128\text{ KiB} = 8 \times (8192 - 1000) \times 128\text{ KiB} \approx 7\text{ GiB}$$

4. What an unmasked empty slot steals ♦

Each zero slot scores 0 and contributes $e^0 = 1$ to the softmax denominator.

$$w_{\text{empty}} = \frac{n_{\text{empty}}}{\sum_i e^{s_i} + n_{\text{empty}}}$$

INTUITION

- ♦ Early in a long reservation almost every slot is empty, so the weight they steal can be huge. Masking to $-\infty$ (or slicing to the pointer) makes it exactly 0.

The four bugs

Bug	Symptom	Fix
RoPE at position 0 in decode	First token fine, then gibberish	Rotate by the pointer's position
Reading empty slots	Attention diluted toward zeros	Slice or mask to the filled length
Not resetting	Next user attends to the previous user's text	Reset the pointer per request
Copying KV heads to 32	◆ 4× the memory (512 KiB/token)	Store 8, broadcast at attention

Three tests

Test	Pass criterion
Correctness	100 greedy tokens identical with/without cache; logits within tolerance (16-bit)
Memory	Allocated bytes = $2 \cdot \text{layers} \cdot \text{KV heads} \cdot \text{head dim} \cdot \text{length} \cdot \text{bytes}$
Speed	Per-step time nearly flat with the cache, climbing without

Rapid-fire Q&A

Question	Crisp answer
Why preallocate instead of torch.cat?	cat copies everything every step (quadratic); preallocation is one write and fixed shapes.
Why do fixed shapes matter?	CUDA graphs and compilers need static shapes.
What does the pointer do?	Marks the filled length: where to write next and how far to read.
Why aren't cached logits bit-identical?	16-bit kernels compute in a different order; they agree within tolerance.
What's the cost of preallocating?	Reserved-but-unused memory caps how many requests fit (paged memory fixes it).

TRAP

- Feeding the decode token at position 0: RoPE needs its true position.
- Assuming zero-filled slots are harmless: $e^0 = 1$ gives them weight.
- Expecting bit-identical logits in 16-bit: compare tokens exactly, logits with a tolerance.